
Averaged Jacobian Regularity Can Misrank Few-Step Flow-Matching Schedules: A Certified Gaussian Counterexample

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Few-step sampling in flow matching and stochastic interpolants requires an inter-
2 polation schedule before any endpoint error is observed. A natural question for
3 schedule design is whether a scalar regularity functional ranks two schedules in
4 the same order as equal-NFE discretization error. Chen et al. propose minimizing
5 A_2 , the time-integrated squared spatial Jacobian norm of a flow-matching marginal
6 ODE, without a proved discretization-error bound. For independent $N(0, 1)$ and
7 $N(0, 4)$, exact regularity integrals prefer trigonometric VP while the Heun method
8 at NFE 8 prefers the linear path in Gaussian W_2 . The inversion is certified by
9 an exact rational Heun product and a nonnegative polynomial in π and $\sqrt{2}$, not
10 by floating-point comparison. The unsigned time average discards signed stage
11 information that a fixed-stage Runge–Kutta method uses. All comparisons use
12 exact Gaussian marginal fields; multimode results are finite deterministic censuses.
13 No learned vector field is used.

14 1 Introduction

15 Few-step sampling commits to a path, a solver, and a small NFE budget before any error is observed.
16 Lipschitz constants enter classical one-step bounds [Hairer et al., 1993]. Chen et al. [2025] propose
17 minimizing A_2 , the time-integrated squared spatial Jacobian norm (their Definition 3.2), as a *selection*
18 criterion. They prove no bound relating A_2 to a discretization error. We test whether the pairwise
19 order that A_2 induces between two interpolant paths is reliable at equal NFE on a flow-matching
20 marginal ODE.

21 We isolate that ranking question on commuting Gaussian interpolants, where the field is affine. For
22 the certified scalar pair, the Gaussian drift and W_2 formula are closed form and the regularity integrals
23 are exact, so inversion is a property of the criterion rather than an estimation artefact. Multimode
24 censuses below are deterministic numerical calculations: $\mathcal{R}$ by adaptive quadrature, and W_2 by
25 exact Gaussian moments propagated through modal Runge–Kutta updates. They are not a symbolic
26 certificate. The leading result is a deterministic scalar Heun example (Section 3). A modal identity
27 then records the signed information that $\mathcal{R}$ discards for the implemented Euler and Heun factors
28 (Section 4). Section 5 splits the ranking question into three comparisons: pairwise linear versus VP,
29 pairwise VP versus the Chen Ex. 3.3 scalar schedule, and a four-path census including the per-mode
30 path, which is not a C^1 shared interpolant when Σ_1 has at least two distinct eigenvalues. Section 6
31 records a complementary Class S construction in which the $\mathcal{R}$ -minimizer is Euler-inexact.

32 Contemporaneous work studies Lipschitz Euler rates [Stéphanovitch, 2026], sharpness-aware
33 grids [Gupta et al., 2026], least-action paths [Du et al., 2026], and instance-aware discretizations [Yuan

et al., 2026]. Those results control other objects. Tao and Choi [2026] give the general strain/vorticity form of the phenomenon; Section 7 records the precise relation.

2 Gaussian interpolants and averaged regularity

Let $X_0 \sim N(\mu_0, \Sigma_0)$ and $X_1 \sim N(\mu_1, \Sigma_1)$ be independent. A shared schedule is a scalar pair (α, σ) . We distinguish three regularity classes.

- Class S: $\alpha, \sigma \in C^1([0, 1])$ with $\alpha(0) = \sigma(1) = 1$ and $\alpha(1) = \sigma(0) = 0$. Linear and trigonometric VP belong to Class S.
- Class Q: a C^1 path of SPD covariances $Q(t)$ with continuous $A = \frac{1}{2}Q'Q^{-1}$. The per-mode schedule $q_i(t) = \lambda_i^t$ is Class Q.
- Class R: $\alpha^2, \sigma^2 \in C^1([0, 1])$, $\alpha^2 \geq 0$, $\sigma^2 \geq 0$, and $\alpha(0)^2 = 1$, $\sigma(0)^2 = 0$, $\alpha(1)^2 = 0$, $\sigma(1)^2 = 1$. Square roots need not be C^1 at the endpoints. Chen Ex. 3.3, $\sigma(t)^2 = (M^t - 1)/(M - 1)$ and $\alpha(t)^2 = (M - M^t)/(M - 1)$ for $M > 0$, $M \neq 1$, is Class R and Class Q, not Class S. The $M = 1$ limit is $\sigma(t)^2 = t$, $\alpha(t)^2 = 1 - t$.

The interpolant $X_t = \alpha X_0 + \sigma X_1$ [Albergo et al., 2025, Lipman et al., 2023] is Gaussian with mean $m_t = \alpha\mu_0 + \sigma\mu_1$ and covariance $Q = \alpha^2\Sigma_0 + \sigma^2\Sigma_1$. The marginal field is affine,

$$b(t, x) = A(t)x + c(t), \quad A = C_t Q^{-1}, \quad c = \dot{m} - Am. \quad (1)$$

The formula $C_t = \dot{\alpha}\Sigma_0 + \dot{\sigma}\Sigma_1$ holds on $(0, 1)$; equivalently $C_t = \frac{1}{2}(\alpha^2)'\Sigma_0 + \frac{1}{2}(\sigma^2)'\Sigma_1$, which extends continuously to $[0, 1]$ for Class R. This is a flow-matching marginal ODE, not a score-based probability-flow ODE in the sense of Song et al. [2021]. If $\Sigma_0 = I$, then Q , Q' , A , and Σ_1 share the eigenbasis of Σ_1 , and mode i obeys $x'_i = a_i(t)x_i + c_i(t)$ with $q_i = \alpha^2 + \lambda_i\sigma^2$ and $a_i = q'_i/(2q_i)$. The covariance drift formula is independent of the mean path $c(t)$. We compare the linear path $(1 - t, t)$ and trigonometric VP $(\cos(\pi t/2), \sin(\pi t/2))$. Following Chen et al. [2025, Definition 3.2], $\mathcal{R}[b] = \int_0^1 \|A(t)\|_2^2 dt$. For state-independent A this coincides with their path expectation. By definition $\mathcal{R}$ ignores $c(t)$. Error is Gaussian W_2 of the endpoint laws [Gelbrich, 1990]. This is neither pathwise trajectory error nor local truncation error. Equal NFE uses $S = B/s$ steps with $s = 1, 2, 4$ for Euler, Heun, and RK4.

Any regular variance path from 1 to $\lambda > 0$ obeys $\int_0^1 a dt = \frac{1}{2} \ln \lambda$. Hence $\mathcal{R} = \int_0^1 a^2 dt = (\int a)^2 + \text{Var}_t(a) \geq (\frac{1}{2} \ln \lambda)^2$, with equality if and only if a is constant. The one-dimensional minimizer is $q(t) = \lambda^t$, $a \equiv \frac{1}{2} \ln \lambda$, which is Chen et al. Example 3.3. Linear and trigonometric VP are C^1 shared interpolants. In several dimensions, the per-mode schedule $q_i(t) = \lambda_i^t$ yields $\mathcal{R} = (\frac{1}{2} \max_i |\ln \lambda_i|)^2$. When Σ_1 has at least two distinct eigenvalues, that path is not a Class S shared interpolant.

3 A one-dimensional Gaussian ranking counterexample

Assumption 1. Endpoints $N(0, 1)$ and $N(0, 4)$ are independent. Schedules are linear and trigonometric VP. The solver is explicit Heun with NFE 8 (four steps of size $1/4$).

Proposition 1. Under the assumption, $\mathcal{R}_{\text{lin}} = 5\pi/8 - 1 > \mathcal{R}_{\text{VP}} = \pi^2/16$, while the Heun factors satisfy $r_{\text{lin}} = 6797469/3559400$ and $0 < r_{\text{VP}} < 187/100$. Hence $W_2^{\text{lin}} < 0.091 < 0.130 < W_2^{\text{VP}}$.

Proof sketch. The linear variance is $q_{\text{lin}}(t) = 1 - 2t + 5t^2$. Completing the square gives the equivalent identities $q_{\text{lin}}(t) = 5[(t - 1/5)^2 + (2/5)^2] = 5(t - 1/5)^2 + 4/5$. The substitution $u = t - 1/5 = (2/5) \tan \theta$ converts $\int a_{\text{lin}}^2$ into $(5/2) \int \sin^2 \theta d\theta$. Limits $\arctan(-1/2)$ and $\arctan 2$ together with $\arctan 2 + \arctan(1/2) = \pi/2$ yield $\mathcal{R}_{\text{lin}} = 5\pi/8 - 1$. The VP integral evaluates to $\pi^2/16$ by a Weierstrass identity. The four linear Heun factors are the rationals $47/52$, $6/5$, $497/370$, and $97/74$, all positive. Exact VP grid values are $a(0) = a(1) = 0$, $a(1/2) = 3\pi/10$, and $a(1/4) = (9 + 15\sqrt{2})\pi/82$, $a(3/4) = (-9 + 15\sqrt{2})\pi/82$. Positivity of $a(3/4)$ is $15^2 \cdot 2 - 9^2 = 369 > 0$. The four-step product lies in $\mathbb{Q}[\pi, \sqrt{2}]$ with nonnegative coefficients, so substituting $\sqrt{2} < 99/70$ and $\pi < 355/113$ is valid and yields $r_{\text{VP}} < 187/100$ by an integer comparison in Appendix A. Interval arithmetic is a cross-check only. $\square$




Figure 1: Certified scalar counterexample for independent $N(0, 1)$ and $N(0, 4)$. (a) Closed-form drifts; the shaded region is $a_{\text{lin}} < 0$; markers are Heun nodes at NFE 8. (b) Exact regularity integrals $5\pi/8 - 1$ and $\pi^2/16$. (c) Gaussian W_2 at Heun NFE 8. Symbolic certificate, not a census.

4 Sign change and transported defects

For any independent linear coupling, $a_{\text{lin}}(0) = -1 < 0 < a_{\text{lin}}(1) = +1$. The VP drift vanishes at the endpoints and does not change sign on $(0, 1)$. Averaging a^2 therefore erases a qualitative distinction. Fixed-stage Runge–Kutta methods use signed stage evaluations that $\mathcal{R}$ discards; in the tested schemes, this distinction contributes to the observed inversion.

Let $e_n = \sqrt{q(t_{n+1})/q(t_n)}$ and let r_n be the numerical factor. The telescoping identity

$$\prod_n r_n - \prod_n e_n = \sum_n (r_n - e_n) \left(\prod_{j < n} r_j \right) \left(\prod_{j > n} e_j \right) \quad (2)$$

is an accounting tool, not the content of the inversion. The content is the leading exact-minus-method coefficient: $\frac{1}{2}(a' + a^2)h^2$ for Euler and $(-a''/12 + a^3/6)h^3$ for the implemented Heun factor. Those expansions involve different derivatives, signs, and stage locations.

5 Four-path census and the inversion region

Source $N(0, I)$. Four centered targets: anisotropic Gaussians with condition number 4 in dimensions 2 and 8 (eigenvalues a geometric sequence from $1/2$ to 2); low-rank covariances $FF^\top + 0.05I$ of rank 2 (seed 271828). Because $\mathcal{R}$ uses the spectral norm, the two anisotropic dimensions share the same $\mathcal{R}$ pair $(0.972296, 0.192270)$, so the grid contains three distinct linear-versus-VP $\mathcal{R}$ comparisons. Each geometry is run under Euler, Heun, and RK4 at NFE $\{8, 16, 32\}$. Headline $\mathcal{R}$ is the continuous integral of $\max_i a_i(t)^2$. Ranking inversions were identified in Phase 3; the counts below are a confirmatory Phase 4 reproduction on configurations frozen before the rerun, not a preregistration that predates discovery.




Figure 2: Three comparisons for $\mathcal{R}$. Each row shows continuous $\mathcal{R}$ (left) and Euler NFE 8 Gaussian W_2 (right). (a) linear versus VP on low-rank $d = 8$; (b) VP versus Chen Ex. 3.3 on anisotropic $d = 8$; (c) four paths on low-rank $d = 8$. Counts: 5 of 12 cells, 9 of 36 blocks, 36 of 36 among four candidates. The last does not imply that a global $\mathcal{R}$ -minimizer minimizes fixed-NFE error.

Table 1: Four schedules: linear, trigonometric VP, shared Chen Ex. 3.3 scalar ($\mathcal{R}_{\text{sc}}$, $M = \lambda_{\text{max}}$), and per-mode log-covariance ($\mathcal{R}_{\text{log}}$). $\mathcal{R}$ is continuous; the last column is the W_2 -minimizer at Heun, NFE 8. Among these four candidate paths, per-mode wins Heun-8 on every geometry. This finite-census observation does not imply that a global $\mathcal{R}$ -minimizer minimizes fixed-NFE error. When Σ_1 has at least two distinct eigenvalues, the per-mode path is not a C^1 shared interpolant. The Gaussian drift and W_2 formula are closed form. Deterministic calculation: $\mathcal{R}$ is evaluated by adaptive quadrature, while W_2 uses exact Gaussian moments propagated through the modal Runge–Kutta updates.

family	d	$\mathcal{R}_{\text{lin}}$	$\mathcal{R}_{\text{VP}}$	$\mathcal{R}_{\text{sc}}$	$\mathcal{R}_{\text{log}}$	Heun-8 W_2 min
anisotropic	2	0.972	0.192	0.173	0.120	log-cov
anisotropic	8	0.972	0.192	0.173	0.120	log-cov
low-rank	2	1.246	0.625	0.668	0.380	log-cov
low-rank	8	2.944	4.731	11.419	2.244	log-cov

98 The per-mode log-covariance path $q_i(t) = \lambda_i^t$ is one of four candidate paths. Among those four paths
99 on the stated 36-block census it has the smallest $\mathcal{R}$ on every geometry and the smallest Gaussian W_2
100 in all 36 blocks (Table 1 reports Heun at NFE 8). This finite-census observation does not imply that
101 a global $\mathcal{R}$ -minimizer minimizes fixed-NFE error. When Σ_1 has at least two distinct eigenvalues,
102 that path is not a C^1 shared interpolant. Linear versus VP invert in 5 of 12 geometry $\times$ solver cells; 4
103 of those 12 invert at every tested NFE. As a pairwise comparison within the shared-schedule class,
104 trigonometric VP versus Chen Ex. 3.3 with $M = \lambda_{\text{max}}$ invert in 9 of 36 blocks and 4 of 12 cells:
105 anisotropic $d = 2$, Euler at NFE 16 and 32; anisotropic $d = 8$, Euler at NFE 8, 16, and 32; low-rank
106 $d = 2$, Euler at NFE 8 and Heun at NFE 8, 16, and 32; low-rank $d = 8$, none. The fourteen inverted
107 linear-versus-VP solver-budget rows are a hierarchical listing of those cells, not a population rate.

Table 2: Strongest centered inversion: low-rank, $d = 8$, Euler, NFE 8. Eigenvalues: 0.05 (six times), 4.64827725629505, 10.675638226546441. Deterministic calculation: $\mathcal{R}$ is evaluated by adaptive quadrature, while W_2 uses exact Gaussian moments propagated through the modal Runge–Kutta updates.

path	$\mathcal{R}$	$\widehat{\mathcal{R}}_{24}$	Gaussian W_2
linear	2.9441044083	2.9476523251	0.8108540111
VP	4.7305438136	4.7295206355	0.4564779075

108 A one-dimensional sweep $\lambda \in [0.05, 100]$ (Figure 3) locates the linear-versus-VP inversions. At $\lambda =$
 109 4, the inversion is present under Heun at every tested budget NFE $\in \{4, 8, 12, 16, 24, 32, 64, 128, 256\}$,
 110 absent under Euler, and mixed under RK4 (absent at NFE 8 and 12). Euler *does* invert at more
 111 extreme λ on the same sweep. The scalar example is therefore solver-dependent, not an artefact of
 112 Heun at budget 8 alone.




Figure 3: Linear versus VP ranking inversion versus target variance λ . Dark cells: $\mathcal{R}$ and W_2 disagree. The dashed line is $\lambda = 4$ (Proposition 1).

113 6 A grid-aware Class S Euler construction

114 The 36 of 36 count is a finite census among four candidates. It does not imply that a global $\mathcal{R}$ -
 115 minimizer minimizes fixed-NFE error. The following construction makes that fence explicit. Fix
 116 $L > 0$ and $N \geq 1$, set $h = 1/N$ and $\varepsilon_N = N(e^{L/N} - 1) - L > 0$, and define $a_0(t) = L$ and
 117 $a_{1,N}(t) = L + \varepsilon_N \cos(2\pi Nt)$. Both exact endpoints of $x' = a(t)x$, $x(0) = 1$ equal e^L . The
 118 constant field a_0 uniquely minimizes $\int_0^1 a^2$ among scalar fields with integral L . The oscillatory field
 119 has strictly larger regularity $\int a_{1,N}^2 = L^2 + \varepsilon_N^2/2$. Left-endpoint Euler is exact for $a_{1,N}$ and has
 120 strictly positive endpoint error for a_0 . The embedding $q_k = \exp(2 \int_0^t a_k)$, $\alpha_k = \sqrt{q_k} \cos(\pi t/2)$,
 121 $\sigma_k = \sqrt{q_k/e^{2L}} \sin(\pi t/2)$ lies in Class S. The field $a_{1,N}$ is chosen after the grid; this is not the
 122 mechanism of Proposition 1.

123 7 Related work

124 Flow matching and interpolants supply the Gaussian family [Lipman et al., 2023, Albergo et al.,
 125 2025, Liu et al., 2023, Tong et al., 2024]. Lipschitz-guided design proposes the selection criterion
 126 tested [Chen et al., 2025]. Solver-specific schedules and exact linear parts are documented in
 127 diffusion practice [Sabour et al., 2024, Karras et al., 2022, Lu et al., 2022, Zhang and Chen, 2023].
 128 Stéphanovitch [2026] prove sharp Lipschitz regularity of canonical flow-matching fields and Euler-
 129 type W_2 rates; that is a bound, not a pairwise ranking of two interpolants at equal NFE.

Tao and Choi [2026] decompose $\nabla v = S + \Omega$ and control Euler error by the logarithmic norm $\mu_2 = \lambda_{\max}(S)$, which is signed, whereas $\mathcal{R}$ integrates $\|A\|_2^2$. In one dimension $\mu_2(a) = a$ and $\|A\|_2 = |a|$. Their Theorem 5 states that McCann displacement interpolation has vanishing material derivative, hence Euler is order 2 in the Eulerian analysis; Remark 7 records the stronger Lagrangian exactness observed on that interpolant. Neither statement applies to the independent-coupling linear or VP path studied here. The present paper supplies a rationally certified Gaussian instance of ranking failure for the unsigned average $\|A\|_2^2$. A counterexample to μ_2 itself is outside our scope. Benton et al. [2024] bound W_2 of a learned flow; that is an error bound, not a pairwise ranking. Consistency models [Song et al., 2023] distill few-step maps.

8 Limitations

The systems are commuting affine Gaussian interpolants, three fixed-step Runge–Kutta methods, and small equal-NFE budgets. No learned neural field is evaluated. Chen et al. prove no discretization-error bound and do not claim a universal pairwise ranking at fixed NFE. As a pairwise comparison within the shared-schedule class, VP versus Chen Ex. 3.3 with $M = \lambda_{\max}$ still invert in 9 of 36 blocks. Among four candidate paths, the per-mode path wins 36 of 36 blocks on this census; that does not imply that a global $\mathcal{R}$ -minimizer minimizes fixed-NFE error. The Gaussian family is a probe in which the drift and W_2 formula are closed form; in a learned field the linear part is not exact, so exponential integrators that solve an affine scalar mode exactly are not a drop-in substitute. Cell counts are hierarchical descriptive summaries, not population rates.

9 Conclusion

Averaged squared Jacobian regularity can misrank equal-NFE Gaussian W_2 already for two C^1 interpolants of independent $N(0, 1)$ and $N(0, 4)$. The Heun calculation is exact on the linear path and rationally enclosed on trigonometric VP. The same split appears as a pairwise comparator (5 of 12 geometry $\times$ solver cells), as a pairwise shared-schedule comparison (9 of 36 blocks), and as a four-candidate census (36 of 36) that does not promote A_2 to a global few-step optimality theorem.

References

- Michael S. Albergo, Nicholas M. Boffi, and Eric Vanden-Eijnden. Stochastic interpolants: A unifying framework for flows and diffusions. *Journal of Machine Learning Research*, 26(209):1–80, 2025.
- Joe Benton, George Deligiannidis, and Arnaud Doucet. Error bounds for flow matching methods. *Transactions on Machine Learning Research*, 2024. arXiv:2305.16860.
- Yifan Chen, Eric Vanden-Eijnden, and Jiawei Xu. Lipschitz-guided design of interpolation schedules in generative models. *arXiv preprint arXiv:2509.01629*, 2025. version 3, 16 May 2026.
- Shukai Du, Junzhe Zhang, and Yiming Li. Lagrangian flow matching: A least-action framework for principled path design. *arXiv preprint arXiv:2605.15419*, 2026.
- Matthias Gelbrich. On a formula for the L^2 Wasserstein metric between measures on Euclidean and Hilbert spaces. *Mathematische Nachrichten*, 147:185–203, 1990.
- Aditi Gupta, Soon Hoe Lim, Annan Yu, and N. Benjamin Erichson. Sharpen your flow: Sharpness-aware sampling for flow matching. *arXiv preprint arXiv:2605.11547*, 2026.
- Ernst Hairer, Syvert P. Nørsett, and Gerhard Wanner. *Solving Ordinary Differential Equations I: Nonstiff Problems*. Springer, 2nd edition, 1993.
- Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2206.00364.
- Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matthew Le. Flow matching for generative modeling. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2210.02747.

- 176 Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and
177 transfer data with rectified flow. In *International Conference on Learning Representations (ICLR)*,
178 2023. arXiv:2209.03003.
- 179 Cheng Lu, Yuhao Zhou, Fan Bao, Jianfei Chen, Chongxuan Li, and Jun Zhu. DPM-solver: A fast
180 ODE solver for diffusion probabilistic model sampling in around 10 steps. In *Advances in Neural
181 Information Processing Systems (NeurIPS)*, 2022. arXiv:2206.00927.
- 182 Amirmojtaba Sabour, Sanja Fidler, and Karsten Kreis. Align your steps: Optimizing sampling
183 schedules in diffusion models. In *International Conference on Machine Learning (ICML)*, 2024.
184 arXiv:2404.14507.
- 185 Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben
186 Poole. Score-based generative modeling through stochastic differential equations. In *International
187 Conference on Learning Representations (ICLR)*, 2021. arXiv:2011.13456.
- 188 Yang Song, Prafulla Dhariwal, Mark Chen, and Ilya Sutskever. Consistency models. In *International
189 Conference on Machine Learning (ICML)*, 2023. arXiv:2303.01469.
- 190 Arthur Stéphanovitch. Lipschitz regularity in flow matching and diffusion models: Sharp sampling
191 rates and functional inequalities. *arXiv preprint arXiv:2604.06065*, 2026.
- 192 Chenxi Tao and Seung-Kyum Choi. On the role of strain and vorticity in numerical integration error
193 for flow matching. *arXiv preprint arXiv:2605.06680*, 2026.
- 194 Alexander Tong, Kilian Fatras, Nikolay Malkin, Guillaume Hugué, Yanlei Zhang, Jarrid Rector-
195 Brooks, Guy Wolf, and Yoshua Bengio. Improving and generalizing flow-based generative
196 models with minibatch optimal transport. *Transactions on Machine Learning Research*, 2024.
197 arXiv:2302.00482.
- 198 Liangyu Yuan, Ruoyu Wang, Tong Zhao, Dingwen Fu, Mingkun Lei, Beier Zhu, and Chi Zhang. Few-
199 step diffusion sampling through instance-aware discretizations. In *Proceedings of the IEEE/CVF
200 Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 35882–35892, 2026.
201 arXiv:2603.17671.
- 202 Qinsheng Zhang and Yongxin Chen. Fast sampling of diffusion models with exponential integrator.
203 In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2204.13902.

204 A Rational VP Heun certificate

205 This appendix records the integer comparison used in Proposition 1. It is part of this PDF, not a
206 separate supplement file.

207 The four-step VP Heun product is the following element $P(\pi, \sqrt{2})$ of $\mathbb{Q}[\pi, \sqrt{2}]$, obtained by
208 expanding the four Heun factors with $h = 1/4$ and reducing $\sqrt{2}^2 = 2$:

$$P(\pi, \sqrt{2}) = 1 + \frac{3}{40}\pi + \frac{15}{164}\pi\sqrt{2} + \frac{78579}{10758400}\pi^2 + \frac{45}{5248}\pi^2\sqrt{2} + \frac{11421}{17213440}\pi^3 + \frac{4671}{17213440}\pi^3\sqrt{2} \\ + \frac{5103}{275415040}\pi^4 + \frac{567}{55083008}\pi^4\sqrt{2} + \frac{243}{2203320320}\pi^5 + \frac{729}{2203320320}\pi^5\sqrt{2} + \frac{729}{176265625600}\pi^6. \quad (3)$$

209 Each of the twelve coefficients is a ratio of nonnegative integers, hence nonnegative, and the constant
210 term is 1. Substituting $\pi < 355/113$ and $\sqrt{2} < 99/70$ is therefore valid and yields

$$P\left(\frac{355}{113}, \frac{99}{70}\right) = \frac{192113353671412470139303}{102753427879939708813312} < \frac{187}{100},$$

211 equivalently

$$19211335367141247013930300 < 19214891013548725548089344.$$

212 Hence $r_{\text{VP}} < 187/100$ and $W_2^{\text{VP}} > 13/100$. The bound $\sqrt{2} < 99/70$ is $99^2 - 2 \cdot 70^2 = 1$;
213 $\pi < 355/113$ follows from Machin’s formula with Leibniz remainders (seven Taylor terms for
214 $\arctan(1/5)$, four for $\arctan(1/239)$). An independent mpmath 1.3.0 interval enclosure is a
215 cross-check, not the proof.