
When Averaged Field Regularity Fails to Rank Few-Step Generative Paths

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Choosing a probability path for few-step generative sampling is often guided by
2 averaged regularity of the velocity field: a path with smaller averaged squared Jaco-
3 bian norm is expected to discretize better at a fixed budget of function evaluations
4 (NFE). We study this question in commuting Gaussian flows, where endpoint laws
5 and Gaussian 2-Wasserstein (W_2) errors are available in closed form and the regu-
6 larity integrand is analytically known; sampling, training, and estimation noise are
7 absent. The averaged-regularity ordering of linear versus variance-preserving paths
8 fails to determine the fixed-NFE W_2 ordering: fourteen inversions identified in a
9 registered benchmark reproduce in a frozen clean-code grid across Euler, Heun, and
10 RK4, survive an 80-digit precision audit, and recur in a separately pre-registered
11 non-centered replication family. An exact modal decomposition attributes the
12 mismatch to solver-stage sampling and signed one-step defects transported to the
13 endpoint; an explicit scalar construction proves the non-implication for fixed-grid
14 Euler. This is a controlled limitation study: it does not refute regularity-guided
15 schedule design, and no learned model is evaluated.

16 1 Introduction

17 Few-step sampling of continuous-time generative models commits to a probability path, a solver, and
18 a small NFE budget before any error is observed. Lipschitz and stability constants enter classical
19 numerical-error bounds [Hairer et al., 1993], and Chen et al. [2025] propose averaged squared drift
20 Lipschitzness as a schedule criterion; bounding an error, however, is weaker than determining which
21 of two paths errs less at an equal small budget. We isolate that question where the marginal flow is
22 affine, the fixed-step endpoint law follows in closed form, and the error is the closed-form Gaussian
23 W_2 [Gelbrich, 1990]; reported values are checked against an 80-digit reference.

24 Contributions.

- 25 • A controlled closed-form W_2 benchmark on commuting Gaussian flows: fourteen inversions
26 identified in a registered benchmark reproduce in a frozen clean-code grid (14 of 36 comparisons,
27 80-digit-verified), and a separately pre-registered non-centered replication family shows the same
28 phenomenon (11 of 18).
- 29 • An exact eigenmode analysis: the endpoint factor error is a transported sum of signed, solver-
30 stage-dependent local defects, which an unsigned time average of Jacobian norms cannot retain.
- 31 • An explicit smooth scalar construction proving that, for every $L > 0$ and grid size N , strictly
32 larger averaged squared Jacobian is compatible with zero fixed-grid Euler endpoint error.

33 Solver dependence of good schedules is not itself novel [Sabour et al., 2024, Karras et al., 2022];
34 the contribution is the exact, controlled non-implication and its modal explanation, bounding what
35 averaged regularity certifies about fixed-budget *rankings* in the tested systems.




Figure 1: **Conceptual schematic, not experimental data.** (a) Averaging collapses two time-dependent fields to two scalars, while a solver samples only stage locations (dots and squares). (b) Local defects are signed: same-sign defects accumulate, alternating defects cancel. (c) Transported to the endpoint, the error ranking can oppose the averaged ranking. This illustrates the accounting of Section 3, not the measured mechanism of any specific Gaussian inversion.

2 Gaussian probability flows and averaged regularity

Let $X_0 \sim N(\mu_0, \Sigma_0)$ and $X_1 \sim N(\mu_1, \Sigma_1)$ be independent and let a scalar schedule $(\alpha(t), \sigma(t))$, with $\alpha(0) = \sigma(1) = 1$ and $\alpha(1) = \sigma(0) = 0$, define the interpolant $X_t = \alpha X_0 + \sigma X_1$ [Albergo et al., 2025, Lipman et al., 2023]. The marginal law is Gaussian with mean $m_t = \alpha\mu_0 + \sigma\mu_1$ and covariance $Q(t) = \alpha^2\Sigma_0 + \sigma^2\Sigma_1$, and the marginal velocity field is affine,

$$b(t, x) = A(t)x + c(t), \quad A(t) = C_t Q(t)^{-1}, \quad c(t) = \dot{m}_t - A(t)m_t, \quad (1)$$

with $C_t = \dot{\alpha}\alpha\Sigma_0 + \dot{\sigma}\sigma\Sigma_1$; centered with $\Sigma_0 = I$ this is $A(t) = \frac{1}{2}Q'Q^{-1}$, $c \equiv 0$. We compare the linear path $(1-t, t)$ and the trigonometric *variance-preserving* (VP) path $(\cos \frac{\pi t}{2}, \sin \frac{\pi t}{2})$, the interpolant analogue of VP diffusion schedules (background: Song et al., 2021). Following Chen et al. [2025], the baseline is $\mathcal{R}[b] = \int_0^1 \|A(t)\|_2^2 dt$, evaluated by trapezoidal quadrature of the analytically known spectral-norm integrand; by its own definition $\mathcal{R}$ ignores $c(t)$. Error is Gaussian W_2 , the closed-form distribution-level metric of this schedule-design literature.

A fixed-step solver applied to (1) induces an affine endpoint map, reconstructed from $d+1$ deterministic probes; pushing the source moments through it and applying the closed-form Gaussian W_2 [Gelbrich, 1990] gives the endpoint error (float64, verified against an 80-digit reference below). With $\Sigma_0 = I$ and $\Sigma_1 = U \text{diag}(\lambda_i) U^\top$ all matrices commute: each mode obeys $x'_i = a_i(t)x_i + c_i(t)$ with $a_i = q'_i/(2q_i)$, $q_i = \alpha^2 + \sigma^2\lambda_i$, and centered, $W_2^2 = \sum_i (|r_i| - \sqrt{\lambda_i})^2$ with r_i the product of the numerical one-step factors of mode i .

3 Why the ranking need not agree: transported signed defects

The exact one-step transition of mode i over $[t_n, t_{n+1}]$ is $e_{i,n} = \sqrt{q_i(t_{n+1})/q_i(t_n)}$, while a Runge-Kutta step multiplies by a factor $r_{i,n}$ built from stage samples of a_i (one for Euler, two for Heun, four for RK4). For smooth a , the leading local defect (exact minus method) is $\frac{1}{2}(a' + a^2)h^2$ for Euler and $(-\frac{1}{12}a'' + \frac{1}{6}a^3)h^3$ for Heun—different derivatives, signs, and stage locations. The endpoint factor error obeys the exact telescoping identity

$$\prod_{n < S} r_{i,n} - \prod_{n < S} e_{i,n} = \sum_{n < S} (r_{i,n} - e_{i,n}) \left(\prod_{j < n} r_{i,j} \right) \left(\prod_{j > n} e_{i,j} \right), \quad (2)$$

with no smallness or positivity assumption. This identifies what $\mathcal{R}$ discards: which times the stages sample, the derivatives of A , the *signs* of local defects, their cancellation, and their transport to the endpoint—so an unsigned average of $\|A\|^2$ need not determine the fixed-NFE ordering.

Table 1: Strongest inversion: low-rank Gaussian, $d = 8$, Euler, NFE 8 (artifact A1).

path	$\mathcal{R}[b]$ (lower = preferred)	Gaussian W_2
linear	2.9476523251	0.8108540111
variance-preserving	4.7295206355	0.4564779075




Figure 2: **Left:** the strongest inversion in both quantities; hatching marks VP. **Right:** W_2 margins of all fourteen inversion blocks on a logarithmic axis (which compresses, never exaggerates, magnitudes); marker shape and color give the W_2 -preferred path, always opposite to the baseline (artifact A1).

4 Controlled benchmark and results

Design and chronology. Source $N(0, I)$; anisotropic (condition number 4) and seeded low-rank ($FF^\top + 0.05I$, rank 2) targets; $d \in \{2, 8\}$; linear and VP paths; Euler, Heun, RK4; NFE $\in \{8, 16, 32\}$ with equal-NFE accounting (steps = NFE/stages): 72 float64 CPU configurations in 36 two-path blocks. A *ranking inversion* is a block where the $\mathcal{R}$ and W_2 orderings strictly disagree. The inversions were first identified in a registered benchmark, then the grid above was frozen and rerun from clean code; post-hoc diagnostics are labeled as such.

Fourteen audited inversions. All 72 configurations reproduce the registered run within 10^{-12} ; 14 of 36 blocks are inversions (Figure 2; audit artifact A1). In the strongest, low-rank at $d = 8$ with Euler at NFE 8 (Table 1), the baseline prefers linear whose W_2 is 78% larger (margin 0.3543761036). All 72 float64 W_2 values agree with an 80-digit mpmath reference to at most 9.71×10^{-10} ; the smallest inversion margin, 1.1189×10^{-5} , exceeds this 11,500-fold (artifact A2).

Solver-path interaction. In the low-rank grid, VP is preferred under Euler and linear under Heun and RK4, at both dimensions and every primary budget (Figure 3); the pattern persists at NFE 64 and 128 and under $\pm 10\%$ target perturbations (artifact A4). Solver-dependent schedule preference is established [Sabour et al., 2024]; the point is that no solver-independent scalar can represent it.

Pre-registered non-centered replication. Before execution we froze one supporting family: source $N(\mu_0, I)$, $\mu_{0,i} = 0.75(-1)^i$; target $N(\mu_1, D)$, $\mu_{1,i} = 1 + 0.25i$, D geometric with condition number 6—so $c(t) \neq 0$ in (1)—with the same paths, solvers, and budgets. Eleven of eighteen blocks are inversions ($\mathcal{R}$, which ignores $c(t)$ by definition, prefers VP with margin 0.7513; W_2 prefers linear), all passing the 80-digit audit (max float64 deviation 2.07×10^{-11} ; smallest inverted margin 1.42×10^{-6} ; artifact A5). The family remains Gaussian, diagonal, and commuting: a single supporting replication for the tested systems, not out-of-distribution evidence. No mixture target supports any conclusion here; mixture evidence failed estimator calibration and was excluded.

5 A minimal non-implication for fixed-grid Euler

Proposition 1. Fix $L > 0$ and an integer $N \geq 1$; let $h = 1/N$ and let forward Euler use the left-endpoint grid $t_n = nh$ on $[0, 1]$. Set $\epsilon_N = N\{e^{L/N} - 1\} - L > 0$ and define $f_k(t, x) = a_k(t)x$




Figure 3: Signed low-rank W_2 margin at equal NFE (artifact A1). Positive values prefer VP; negative values prefer linear; the preferred path flips with the solver. The axis is symmetric-log (linear threshold 10^{-4}) since margins span four orders of magnitude and change sign.

89 with $a_0(t) = L$ and $a_{1,N}(t) = L + \epsilon_N \cos(2\pi Nt)$. For $x(0) = 1$: (i) both exact endpoints equal
90 e^L ; (ii) $\int_0^1 a_{1,N}^2 = L^2 + \epsilon_N^2/2 > \int_0^1 a_0^2$; (iii) Euler is exact for $a_{1,N}$; (iv) Euler has strictly positive
91 endpoint error for a_0 . Hence larger averaged squared Jacobian does not imply larger fixed-grid
92 Euler endpoint error.

93 *Proof.* (i) $\int_0^1 \cos(2\pi Nt) dt = 0$, so both coefficients integrate to L and $x(1) = e^L$. (ii) Orthogonality
94 gives $\int_0^1 a_{1,N}^2 = L^2 + \epsilon_N^2/2$, and $\epsilon_N > 0$ since $e^z > 1 + z$ for $z > 0$. (iii) $\cos(2\pi Nt_n) = 1$ at every
95 node, so each Euler factor is $1 + h(L + \epsilon_N) = e^{L/N}$; the product is e^L . (iv) $(1 + L/N)^N < e^L$ since
96 $\log(1 + z) < z$. $\square$

97 The quantifiers matter: $a_{1,N}$ is chosen after the grid, its amplitude vanishes as $N \rightarrow \infty$, and shifted
98 grids or multi-stage solvers sample different phases. The construction shows non-implication for
99 fixed-grid left-endpoint Euler only—neither the mechanism behind every Gaussian inversion nor a
100 new numerical-analysis phenomenon; the full proof and edge-case audit are in the supplement.

101 6 Related work

102 Flow matching and stochastic interpolants define the Gaussian path family and its fields [Lipman
103 et al., 2023, Albergo et al., 2025, Liu et al., 2023]; Lipschitz-guided schedule design proposes the
104 criterion we test [Chen et al., 2025]. Fast diffusion solvers exploit exactly integrated linear parts
105 and solver-specific expansions [Lu et al., 2022a,b, Zhang and Chen, 2023], EDM separates schedule
106 from discretization [Karras et al., 2022], and Align Your Steps tunes schedules per solver, model,
107 and dataset [Sabour et al., 2024]: solver-dependent schedule quality is known. Runge–Kutta and
108 backward error analysis supply the local-error machinery [Butcher, 1963, Hairer et al., 1993, 2006];
109 Gaussian W_2 is due to Gelbrich [1990]. None states the commuting inversion benchmark or the
110 grid-aware non-implication; search absence does not establish novelty.

111 7 Limitations and conclusion

112 Our systems are closed-form, commuting, affine Gaussian flows with two paths, three solvers,
113 $d \in \{2, 8\}$, and small budgets, plus one pre-registered non-centered replication family; no learned
114 neural field is evaluated, no image, video, or world-model improvement is demonstrated or implied,
115 and the proposition is grid-aware. Exponential and exact-linear integrators solve affine fields exactly:
116 the result concerns generic fixed-stage Runge–Kutta discretization, not best-possible solvers; learned
117 fields are generally non-affine, but no learned-field result is established here. We rank two fixed paths
118 at the identity time parameterization; transfer-based schedule optimization [Chen et al., 2025] is
119 complementary and not addressed. Averaged regularity may remain valuable for schedule design, and
120 our result does not refute Lipschitz-guided schedules: in the tested systems the averaged-regularity
121 ordering does not determine the fixed-NFE W_2 ordering, because fixed-budget error lives in solver-
122 stage samples and transported signed defects the average cannot see. All numerical statements resolve
123 to checksummed audit artifacts (A1–A5) in the supplement.

References

- Michael S. Albergo, Nicholas M. Boffi, and Eric Vanden-Eijnden. Stochastic interpolants: A unifying framework for flows and diffusions. *Journal of Machine Learning Research*, 26(209):1–80, 2025.
- John C. Butcher. Coefficients for the study of Runge-Kutta integration processes. *Journal of the Australian Mathematical Society*, 3(2):185–201, 1963.
- Yifan Chen, Eric Vanden-Eijnden, and Jiawei Xu. Lipschitz-guided design of interpolation schedules in generative models. *arXiv preprint arXiv:2509.01629*, 2025.
- Matthias Gelbrich. On a formula for the L^2 Wasserstein metric between measures on Euclidean and Hilbert spaces. *Mathematische Nachrichten*, 147:185–203, 1990.
- Ernst Hairer, Syvert P. Nørsett, and Gerhard Wanner. *Solving Ordinary Differential Equations I: Nonstiff Problems*. Springer, 2nd edition, 1993.
- Ernst Hairer, Christian Lubich, and Gerhard Wanner. *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*. Springer, 2nd edition, 2006.
- Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2206.00364.
- Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matthew Le. Flow matching for generative modeling. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2210.02747.
- Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2209.03003.
- Cheng Lu, Yuhao Zhou, Fan Bao, Jianfei Chen, Chongxuan Li, and Jun Zhu. DPM-solver: A fast ODE solver for diffusion probabilistic model sampling in around 10 steps. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022a. arXiv:2206.00927.
- Cheng Lu, Yuhao Zhou, Fan Bao, Jianfei Chen, Chongxuan Li, and Jun Zhu. DPM-solver++: Fast solver for guided sampling of diffusion probabilistic models. *arXiv preprint arXiv:2211.01095*, 2022b.
- Amirmojtaba Sabour, Sanja Fidler, and Karsten Kreis. Align your steps: Optimizing sampling schedules in diffusion models. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2404.14507.
- Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations (ICLR)*, 2021. arXiv:2011.13456.
- Qinsheng Zhang and Yongxin Chen. Fast sampling of diffusion models with exponential integrator. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2204.13902.